

Plurimo

Business Plan & Vision Document — Infrastructure for the Independent Knowledge Economy

Note: this business plan is included to show what I plan to do with any money I can make from selling RICD collapsal-causal trackers, as well as to further demonstrate why I want to test RICD against other systems before using it as the theory that powers Plurimo's algorithm.

Executive Summary

Plurimo is a democratized free-thought platform beginning as teaching-optimization software for private tutors and independent educators and scaling into infrastructure for the global independent knowledge economy.

The name is a portmanteau: Pluri, from *e pluribus unum* — out of many, one — fused with Mo, from Mohini, Vishnu's enchantress avatar who held the demons spellbound with her dance while distributing the amrita, the elixir of immortality, to the gods; this represents Plurimo finding its way around institutional gatekeepers to give knowledge—and the earning power it confers—back to the people.

At its core, Plurimo trains AI to optimize human intelligence for systems thinking and cross-domain expertise, rather than replacing human thinking outright. If successful, the AI trained on this platform stands to substantively accelerate human capability rather than substitute for it.

The tutoring platform is the online bookstore that Amazon started as. What Plurimo will become is where Amazon ended up.

Everything in this document beyond the tutoring platform itself — the academic forum, the social layer, the family of domain-specific AI engines, the Knowledge Shepherd economy — is the destination, not the starting point. Plurimo begins as one thing, built and proven completely before the next is added, in a specific, deliberate order: the tutoring platform and a modest social forum launch together first, with the academic forum, the classifieds, and the full-scale social forum arriving later, together, once the domain-specific AI engines reach market and begin making rote education itself obsolete. Each layer is real infrastructure in its own right, not a teaser for what comes after; but the sequence matters, and this document is structured to make that sequence impossible to miss.

The Founder

Plurimo is founded by a working scientist and complexity theorist, not solely a pedagogy expert who happened to also build software. Over roughly two and a half months, the founder independently developed Recursive Information-Container Dynamics (RICD), a formal mathematical framework for adaptive-system collapse and health, validated against real institutional data with a 100% cross-validated accuracy result on a 54-institution panel. RICD is the mathematical engine behind Plurimo's own anti-capture and site-health algorithms — the

same underlying architecture that identifies institutional decay is directly applicable to identifying epistemic capture, forced semantic drift, and community health on the platform itself.

Separately, the founder is also the creator of the Web-Field Model of Cognitive Trait Distributions and Social Evolution, which includes the Cognitive Trait Manifold Ratio (CMTR) theory, the direct conceptual basis for the pedagogical (learning-style) evaluation every teacher and student takes before using Plurimo, and for the platform's own teacher-student-preset matching engine described later in this document.

The founder personally originated the preset system itself and is directly coding the first Writing presets for the Writing Intelligence Engine. Venture funding is what allows this to scale beyond one person, sourcing additional master teachers to work alongside coders and LLM tooling to build out presets in every other subject.

The founder's own track record substantiates this directly rather than as an unverified claim. A long-time teacher at academies, universities and colleges, the founder observed that even at the larger, time-limited class level, it was possible to unlock student talent within the writing and literary domain by appealing to their existing knowledge base, and then across five years of private academic coaching, approximately 90% of middle and high school students she worked with long-term reached placements in major national and international competitions within one to two years of beginning private coaching, including such contests as the Scholastic Art & Writing Awards, National History Day, the Bow Seat Ocean Awareness Contest, the Berkeley Economic Review High School Essay Contest, and the John Locke Institute Global Essay Prize. Students went on to admission in fields as far apart as cognitive neuroscience and international nursing, at institutions including UC Berkeley and the University of Chicago, itself a direct demonstration of the cross-domain mastery this method is built to produce, not a separate credentialing point. Separately, as an educational therapist working one-on-one with a student with autism, ADHD, and dyslexia, the same underlying method increased literacy and reading comprehension by over 80%. The founder holds an MFA in Creative Writing from Hollins University (Summa Cum Laude, 4.0 GPA) and a BA in English with a Spanish minor from Northern Arizona University (Summa Cum Laude, 3.7 GPA).

The Founder's Pedagogical Method: Top-Down, Not Bottom-Up

The founder's own teaching record demonstrates a specific, repeatable result: using a student's own CMTR (their cognitive and learning-style profile) to teach them how to apply what they already know in one domain to unlock rapid mastery in an entirely different one.

The founder asserts that the 10,000-hour rule for mastery of a subject is correct but incomplete. Even children, regardless of formal training, have already accumulated well over 10,000 hours of living inside and observing real systems by the time they reach adolescence. Taught explicitly how to recognize and cross-apply the structural patterns they already understand from their own areas of genuine interest, people—both children and adults alike—require dramatically fewer additional hours to reach real mastery in a new domain. This is top-down education, starting from the structure a student already holds and mapping it onto new material, rather than bottom-up rote accumulation from zero.

Function Before Form, First Principles, Macro to Micro

The founder's teaching method holds a single governing principle throughout: form is subservient to function, and function is taught first. A definitional argument, for example, is never introduced as a formula. Students first cover discourse community, formal logic including logical fallacies, audience, voice, and an introduction to relevant philosophy — the actual functional apparatus a definitional argument exists to serve — before the form of the essay itself ever comes up.

Rather than handing students a structure like the five-paragraph essay, students are taught to look for the gaps — something ill-defined or never defined at all — decide for themselves what they are defining, and build a custom container suited to their own research question. Research itself is taught as an ongoing task, not a discrete unit to complete once: students learn to place source summaries directly beneath the paragraphs where they'll be used, working toward synthesizing sources in their own heads without the summary text in front of them; and once a student is working to answer a real question they set for themselves — rather than completing a task or filling in a formula — research stops being a chore they endure and instead becomes something they look forward to (or, quite often, an all-consuming obsession until the answer is found). Since the idea belongs to the student as does the resulting final outcome—whether that be an essay or a formula—the student becomes quite invested in the process of perfecting and refining it, which is how to ignite a passion for scholarship in any student: by giving them full ownership of their ideas.

This is a consistent method, not a collection of separate techniques: function before form, first principles before structure, macro before micro, in every subject and every level. The founder's own repeated result, borne out across a fifteen-year teaching record: polymathic thinking is the human norm, not a rare exception, and is drilled out of people by educational and career siloing — by telling people what they can and cannot do, rather than showing them how their own particular form of intelligence already applies to the domain in front of them. Polymaths are not mysterious geniuses who simply “learn things immediately; they simply have innate gifts for pattern recognition across domains, which allows them to master crafts and disciplines with relative ease and substantively contribute across disciplinary lines. Plurimo's method exists to make that same capacity accessible to anyone, not only those whose polymathic capabilities have proved unsuppressable in a pedagogically and cognitively oppressive society.

The Problem

Private Tutors and Independent Educators Are Underserved

The private tutoring market is large, growing, and almost entirely without purpose-built software. Existing tools fall into two inadequate categories: full LMS platforms (Canvas, Blackboard, Moodle), designed for institutions and far too complex and expensive for independent tutors; and generic lesson planners, with no pedagogical depth, subject-matter intelligence, or community.

The result is independent tutors cobbling together Google Docs, email, Venmo, and generic scheduling tools — a fragmented, unprofessional experience for both teacher and student.

Beginning Teachers Lack Access to Expert Pedagogical Knowledge

A beginning teacher — a new private tutor, a homeschooling parent, an undertrained classroom teacher — has no reliable mechanism for accessing the accumulated pedagogical wisdom of master teachers in their subject. That knowledge lives in individual experts and largely dies with them or stays locked in expensive professional development programs.

The Intellectual Honesty Problem

Across educational contexts — from homeschools to university classrooms — there is a widespread and growing concern that teaching presents conclusions rather than cultivating thinking. A single theoretical framework is taught as the whole of a subject. Students graduate without understanding that serious, intelligent people disagree about most important questions. No existing platform addresses this structurally. Plurimo does.

Academic Isolation

Independent researchers and academics working in niche fields have almost no infrastructure for finding peers outside of expensive conferences or institutional networks they may not have access to. Important connections between minds working on related problems never happen because there is no forum for them.

PHASE ONE (THE MVP): The Tutoring Platform

This is what gets built and shipped first. Nothing else in this document is required for launch.

Teaching-Optimization Software, Not an LMS

Plurimo is not a course management system. It is teaching-optimization software — a fundamentally different category. Rather than managing content delivery, it optimizes the act of teaching itself through an intelligent preset system.

The Preset System

The heart of the platform is the preset system: structured lesson frameworks designed by expert teachers in each subject. A preset does not merely organize materials — it instructs the teacher in the flow of the lesson and actively facilitates delivery. This makes Plurimo an expertise transfer

mechanism: the pedagogical judgment of a master teacher is embedded in a structure that a novice can follow and learn from simultaneously.

As beginning teachers use the presets repeatedly, they absorb the logic behind good lesson flow. The tool trains them while they use it — a compound effect that improves teacher quality at scale with no additional intervention required.

Presets are built at the level of the individual lesson, not the generic subject. There is no single, undifferentiated “rhetoric and composition” preset — there is a definitional-argument preset, a compare-and-contrast preset, a narrative-essay preset, each its own distinct structure, because each actually requires a different lesson flow to teach well.

Presets range from genuinely useful to highly specific, becoming more useful as the range is traversed rather than merely unlocking more of the same thing. This range applies across every format Plurimo supports — the individual 1:1 lesson is only one example among several, alongside a classroom format, a survey-course format for very large classes, a conference format, and others — running from genuinely useful, subject-specific generic presets at the base (a generic essay preset for the individual-lesson format, for instance, walking through discourse community, formal logic and logical fallacies, audience, voice, a basic introduction to relevant philosophy, an introduction to the thesis concept with a thesis workshop, an introduction to research with a research workshop, an introduction to brainstorming with a brainstorming session, and building the student's own custom container, with the other formats reaching a comparable level of development), up to the fully specific individual essay type within the individual-lesson format at the top of the range: definitional argument, argumentative essay in each of the three classical rhetorical modes (classical, Rogerian, Toulmin), argumentative research paper, analytical research paper, exploratory research paper, compare and contrast, cause and effect, division and classification, textual analysis, visual analysis, narrative essay, and others, each its own dedicated preset — with the classroom, survey-course, and conference formats reaching the same degree of highly specific development. There is no separate, wholly-empty placeholder tier below this — since free and basic share identical preset access, the useful, subject-specific generic preset is the actual floor, not an empty shape a teacher must populate entirely from scratch.

Free and basic tiers have identical access to this entire range of presets, top to bottom — a free-tier user is not limited to the generic, placeholder end of it. Basic adds three structural capabilities, not more presets: for teachers, searchability in the directory, meaning a basic-tier teacher can actually gain new students through the site itself — solving what is, without real competition, the single hardest recurring problem an independent tutor faces: finding new students and establishing credibility with them without a boutique agency's name behind them. A modest monthly fee that directly solves the hardest part of running an independent tutoring practice is not a marginal upgrade; it is the actual reason to pay at all; for students, the ability to find and browse teachers through the site, rather than only bringing an existing tutor with them; and for both, access to the academic and social forums, plus the ability to leave and receive reviews. Reviews directly affect a teacher's own placement in search results; a teacher with a consistent record of good reviews is bumped higher, making review quality a real, ongoing input to visibility, not just a static credential. A free-tier user gets access to the complete basic preset library but is not searchable, cannot browse for teachers, cannot use either forum, and does not participate in the review system. Without retaining the basic presets, the free tier would offer

little beyond what a teacher bringing their own existing students could do on their own, while the forums, reviews, and searchability together are what make the paid tier a genuine business tool rather than only a more polished version of the same private arrangement.

Most free-tier users are expected to elect to train the CMTR site-matching algorithm in exchange for free premium-tier access, per the Site Algorithm Training Track described later in this document. A real minority — small schools and similar institutional users, in particular — will not take that option and will simply remain on the free tier as described above.

Every preset is built from individual cells — discrete, self-contained units that can be edited, reordered, or deleted freely, making every preset fully customizable at every tier, rather than a fixed sequence a teacher is locked into.

All tiers include preloaded teaching materials for the underlying concepts every essay type draws on — e.g. discourse community, logical fallacies, and others — delivered through text and video paired with rapid pattern-matching exercises: a student reads a passage or watches a short video and identifies the discourse community it belongs to from a multiple-choice set, then is presented with the same concept out of its original context — a lawyer speaking on a talk show rather than in a courtroom, two doctors discussing a case in a café rather than a hospital — and asked to identify the discourse community again, to confirm the pattern-recognition genuinely transfers rather than merely matching surface context. Logical fallacies are taught the same way, in both text and video-clip form, with the same rapid pattern-matching structure.

What actually distinguishes premium from free and basic is not this underlying conceptual material, which all tiers include, but two further things. First, completeness of the specific, essay-type-level content: at free and basic, a teacher gets the shared conceptual materials but must upload their own materials specific to the particular essay type being taught, while at premium, everything is entirely pre-planned — every material and every lesson plan for every listed essay type — while remaining fully and easily customizable at the cell level regardless. Second, premium alone offers multiple distinct versions of the same underlying preset, built for different student or teacher Cognitive Trait Manifold Ratio profiles: a teacher can try different versions to see which fits a given student, or open several versions of the same preset simultaneously and drag and drop individual cells between them, mixing pedagogical styles and tools directly rather than choosing only one version wholesale.

A separate integrity mechanism applies to preset content at every tier: if a teacher enters content that is a blatant mismatch for the course or the preset it sits in, that content is automatically flagged and frozen — inaccessible to students — until it is removed. A teacher who believes the flag was wrong may appeal to still use the content, and a human moderator reviews and resolves the appeal within 24 hours.

The CMTR Pedagogical Evaluation

All teachers and all students are required to take a pedagogical evaluation before using the site, grounded directly in the founder's own Cognitive Trait Manifold Ratio theory — establishing each user's actual cognitive and learning-style profile as real, structured data from day one, rather than as a marketing feature added later. This evaluation is what later powers accurate teacher-student-preset matching as the platform scales.

The Classification and Flag System

Every course on Plurimo must be accurately classified to prevent epistemic capture. The platform's intelligent flagging system monitors whether course content matches its stated classification and automatically updates the course heading if it does not.

For example, a course covering only Western History cannot call itself “World History.” It must call itself “Western Theory” or “Occidental History.” A course called “World History” must include Eastern as well as Western History, African history, Oceanian history, and the history of the indigenous peoples of the New World as well as every continent, as well as an outline of the global scholarship surrounding prehistory, to earn itself the title World History; if it does not include all of these, the system flags the mismatch and reclassifies accordingly.

This is not content moderation. It is honest labeling, enforced structurally, without human moderators, without ideological interference, and without restricting what any teacher can teach. Teachers retain complete freedom over content. They cannot misrepresent its scope.

The result: course listings on Plurimo are trustworthy in a way no other platform can claim.

Worth being explicit about the relative technical bar here, since it is easy to overstate: automated course-level classification is comparatively simple pattern-matching against a stated syllabus, not a heavy technical lift in its own right. It is minor compared to what this plan already commits to building regardless — the Writing Intelligence Engine's thousands of preloaded, pre-coded error tags, described later in this document, are a substantially heavier undertaking, and course classification rides on infrastructure the platform needs to build anyway rather than requiring its own separate, comparably difficult engineering effort.

Master Teacher Approval: Results, Not Credentials

Master-teacher status is granted on the basis of demonstrated results, not formal qualifications. Because pedagogical mastery can be demonstrated in genuinely different ways across subjects and teaching contexts, there is no single required credential.

Every portfolio submitted for master-teacher review is evaluated by a human moderator — never by an automated check alone — and that moderator must themselves be a master teacher hired specifically to evaluate applicants in their own subject area. A mathematics pedagogical portfolio is reviewed by an approved master math teacher; a writing pedagogical portfolio by an approved master writing teacher.

Business Model, Phase One

- Free — full platform access, basic lesson planning, forum access, free access for small schools and hybrid courses.
- Basic (Paid) — searchable in the teacher directory, student reviews enabled, generalized course-type presets.

- Premium (Paid) — highly specialized, subject-specific presets designed by verified expert teachers, priority placement.
- Expert Verified (Earned, Not Purchased) — granted via results-based portfolio review; free access to all premium presets; can train intelligence engine in exchange for equity in the resulting software.

Market & Users, Phase One

- Private tutors/academic coaches — the initial primary user base; a large, fragmented market with no purpose-built software.
- Homeschooling families — a large and rapidly growing community; homeschooling parents are often teaching without pedagogical training, making the preset system especially valuable.
- Small schools — institutions that cannot afford or do not need a full LMS; the free tier serves them completely.

The global private tutoring market is valued at over \$200 billion and growing. The homeschooling market in the US alone serves over 3.3 million students.

The Social Forum (Modest Launch)

A third forum launches alongside the tutoring platform itself, kept deliberately small in scope at this stage: an after-class hangout space, open-ended and Reddit-like in structure, with two firm content restrictions — no pornographic content, and no content that promotes, revels in, or treats real-world violence as entertainment; video games lacking any real-world physical component are also not permitted. This mirrors, at a smaller scale, the same discipline governing the tutoring platform's own launch: prove a modest version works before the scope grows.

Content quality within this forum uses the same judgment applied elsewhere on the platform: abusive content is removed outright, low-effort non-abusive content is pushed down in notifications, and high-effort content is surfaced prominently. Human moderators are required at first, training the AI that will eventually take over this judgment as well.

The Phase Transition: Why Phase Two Is Inevitable, Not Optional

Phase Two is not an optional growth lever layered on top of a successful Phase One. It is the platform's necessary response to a structural fact built into Phase One's own design: by building the best possible tutoring-optimization platform, Plurimo is simultaneously compiling the world's most sophisticated, ethically sourced pedagogical training dataset. The moment the domain-specific intelligence engines trained on that dataset — beginning with the Writing Intelligence Engine, continuing through the STEM, Humanities, and Arts engines described later in this document — become mature enough to handle rote instruction directly, they inherently commoditize the original human tutoring marketplace they were trained from.

This is a real, built-in expiration date on Phase One's own business model, not a hypothetical risk. Reaching that point takes years, not months; the engines have to actually mature to the point of handling rote instruction credibly, which requires sustained data collection across a large base

of active master teachers first. But if Plurimo is successful, the transition is a when, not an if, and treating it as a distant hypothetical rather than a structural certainty would be dishonest about what the platform's own core product is actually building toward.

If Plurimo does not build the Academic Forum, the Classifieds, and the Knowledge Shepherd economy ahead of that transition, it builds an engine that destroys its own user base once it succeeds. Building them so that they arrive precisely as the engines mature transforms the transition into a managed handoff instead: displaced human educators move into Knowledge Shepherd roles exactly as the AI automates the mechanics they used to teach directly, with a real, already-built destination waiting for them rather than an afterthought scrambled together once the disruption has already happened. This is the platform's own creative-destruction loop, made a deliberate feature of the design rather than an unaddressed risk discovered later.

PHASE TWO: The Academic Forum, the Classifieds, and the Knowledge Shepherd Economy

Rolls out once the domain-specific intelligence engines reach market and begin displacing rote tutoring demand — years after Phase One's launch, timed to the transition described above, not to a fixed calendar date.

Knowledge Shepherd Presets

As the engines mature, the preset system's own objective function shifts. Phase One's presets exist to transfer subject-matter mastery — teaching an essay type, a math concept, a coding language. Knowledge Shepherd Presets exist for a different purpose entirely: guiding a Shepherd and a polymathic student through the actual process of mapping one mental model onto another, to spot and act on emergent problems no single existing field already owns. The shift is from subject-matter acquisition to cross-domain application, and the presets themselves are built for that shift specifically, not adapted from the Phase One preset library.

The Three Responsibilities of the Knowledge Shepherd

A polymathic student's real obstacles are rarely raw intelligence. They are translation, uneven skill profiles, and isolated execution; a student with a genuine breakthrough insight in one field can lose months trying to learn an entire adjacent field from scratch just to apply it, or stall out because one specific prerequisite skill is missing. The Knowledge Shepherd exists to remove exactly these three obstacles.

Cross-Domain Translation: the Shepherd uses Knowledge Shepherd Presets to map the structural logic of a student's native domain directly onto a target domain, isolating only the small number of foundational principles actually needed to make the cross-application work, bypassing the institutional gatekeeping of full coursework and formal prerequisites entirely.

Targeted Shore-Up: polymathic thinkers are frequently asynchronous learners, capable of graduate-level systems thinking while genuinely missing a basic, specific technical skill the application requires. The Shepherd uses the platform's own modular cell architecture to inject exactly that missing competency, hyper-focused and no broader than the gap actually is.

Field Harmonization: the highest-leverage responsibility, and the one already described elsewhere in this document as the Shepherd's central function: continuously scanning the academic and social forums for complementary minds working on adjacent, unconnected problems, and actively introducing them to see whether a new field capable of solving a real problem emerges.

The Academic Forum

Plurimo includes an open, annotatable database of academic work. Users can save annotations on papers and studies, and converse with each other about specific works in subject-matter forums.

Academic papers become living documents rather than static publications. A layer of collective intelligence — from credentialed academics and serious amateurs alike — accumulates on top of existing knowledge. The serious amateur—the independent historian, the citizen scientist, the self-taught philosopher — finds a home and a community.

Comment quality on the academic forum is governed by its own dedicated mechanism: abusive content is flagged and removed outright. Low-effort, non-abusive content is immediately routed into an adjacent zip file attached to the forum — preserved, not deleted, but no longer visible on the paper itself. Only high-effort, non-abusive comments remain directly on the paper, where others can instantly read them, agree or disagree, and add their own responses. Human moderators are required for this at first, and are themselves training the AI that will eventually take over the classification judgment, the same pattern used elsewhere on the platform for tag-fidelity and work/play-fidelity moderation.

Papers and projects regularly spawn new papers and projects in the following manner: if two or more people comment on another user's comment on the same paper or project specifically to flag an unresolved question, another forum opens up automatically, and places the commenters—both the initial commenters and those who commented on their comment/on the same area of content with a similar or related comment—into a new forum with the original paper or project creator, to create a new paper or project to answer the question; if the original creator is not interested, he or she is free to leave the forum, but the commenters are likewise free to continue their own project on the question, should they wish. If everyone leaves the forum, the folder is closed automatically and any contents saved to a zip file in the original forum where the original paper or project under discussion was uploaded.

In addition to users being able to upload content to create new folders in the academic forum, and folders opening automatically that match commenters with creators, Knowledge Shepherds regularly create forums and place two or more users within them, with instructions to work on a project answering a real-world problem; Knowledge Shepherds retain royalties for a limited number of years from any applications which come out of such ventures, as creation of new fields is a part of their job description.

The content of the individual folders in the academic forum is moderated at first by humans who train AI replacements, similarly to content moderation: anything abusive or inappropriate is flagged and removed; low-effort since work is placed in an archive at the bottom of the forum, to preserve forum history; and high effort work, which cites sources and is academically defensible,

is allowed to stand, regardless of any controversy surrounding it or incompleteness of the project. The rules for ensuring content visibility (editing for logic, formatting, and citations) are clearly stated on the academic forum's main page.

How the Layers Interlock

At this phase, four layers of the platform operate as one synchronized system rather than as separate features. The Academic Forum supplies the raw material: the papers, theories, and open problems a Shepherd draws on. Knowledge Shepherd Presets supply the cognitive tools for actually cross-analyzing that material against a given student's own domain. The Classifieds supply the economic destination, where a resulting breakthrough becomes real, fundable, paid work. And the Social Forum, having scaled beyond its modest Phase One launch into a full generative space by this point, supplies a lower-stakes environment where accidental cross-discipline connections can be tested informally before they become a deliberate Shepherd-facilitated pairing worth the Academic Forum's higher-stakes commitment.

The Forced Semantic Drift Flag and Warning System

Launching at this phase, alongside the academic forum itself rather than as part of Phase One's leaner initial build: a separate mechanism guarding against epistemic capture through forced semantic drift, i.e. coordinated attempts to redefine a term's meaning across the tutoring platform's course classifications or the academic forum's discourse, whatever the stated justification, in order for a subfield or ideology to capture an entire field or a subject area. This applies specifically to the tutoring and academic sides of the platform and deliberately does not extend to the social forum, where apart from abusive language users retain full freedom of speech, including the freedom to be wrong or to lie. Detection is algorithmic, built on RICD's own site-health architecture, tracking rate of linguistic or classificatory change and organized opposition to or suppression of that change as its core signals; an entire team, including linguists and political scientists alongside engineers, is dedicated to building and refining this specific algorithm, not engineers alone. Enforcement is a warning-then-ban structure: a user flagged for a specific instance of forced drift receives a warning to stop that particular instance. Persisting in that same instance after the warning results in a ban. Being caught up in a new, unrelated instance later — even repeatedly, even weekly — does not compound against a user who stopped each time they were warned; only continued persistence within one specific instance escalates to removal.

The forums serve multiple functions at once: subject-matter discussion across discourse communities and disciplines; organic tutor discovery, since a user who finds a compelling discussion can see the teacher's profile and hire them directly; connection infrastructure for independent academics in niche fields who cannot find peers through institutional channels; and a reputation and credibility layer for independent knowledge workers. Forums are free to all users to maximize traffic, community growth, and the network effects that make the platform increasingly valuable over time.

Together, the teaching layer built in Phase One and the research layer built here — connected by community — begin to constitute the full architecture of the independent knowledge economy.

Expansion Markets Unlocked at This Phase

- Independent academics and researchers, for webinars, seminars, and finding peers in niche fields globally.
- Authors, artists, and creative professionals: domain experts who want to teach workshops or masterclasses but lack pedagogical infrastructure.
- Corporate training departments: subject matter experts who are not natural teachers.
- Developing-world educators: where qualified teachers are scarce and the preset system provides expert guidance at no cost.

A pairing that survives and produces a genuine application earns the originating Shepherd ongoing royalties, for as long as that application generates value, structurally similar to an author's reprint royalties or an actor's residuals.

Four open questions are named directly rather than smoothed over: attribution once a field outgrows its founding pair; friction against over-proposing low-quality pairings; whether royalty rights require upfront creator agreement; and the legal and financial structuring needed to keep 'royalties for life' genuinely lucrative while remaining bounded to Plurimo's own actual capacity. On the last point specifically: an uncapped, permanent royalty is very likely the wrong final shape, given the platform's own exposure compounds with every additional field that succeeds over time. The more defensible direction is a capped percentage of the value a founded field generates, with a stated sunset period rather than a truly indefinite term; the exact percentage and sunset length are financial and legal decisions still to be made, not asserted here as settled.

The royalty itself, within whatever cap and sunset structure is ultimately set, is non-negotiable, not a cost center to be minimized. A capped royalty is a legal and financial liability worth managing carefully, but a royalty removed entirely would remove the actual incentive that makes Knowledge Shepherding worth pursuing as a career at all, rather than an ordinary, replaceable gig. The cap protects the platform's own financial exposure; the royalty's existence at all is what protects the talent from exploitation and makes the role a genuinely sought-after position rather than another gig economy race to the bottom.

The Classifieds: The End of the Single-Track Career

A dedicated jobs sub-forum within the academic forum, functioning as an open classifieds, lists newly opened fields, the problem being solved, the amount being offered, and the skills required. This is open to anyone who comes onto the site, not restricted to registered Plurimo users, since the goal is matching the right person to the problem regardless of where that person is coming from.

The classifieds is not a minor feature bolted onto the Knowledge Shepherd economy. It is the mechanism by which the platform's underlying thesis becomes an actual career structure: that single-track careers, and the disciplinary siloing they enforce, are not a natural or necessary feature of how skilled work has to be organized, but simply the logical consequence when there is no infrastructure connecting a person's underlying capability to a newly emerged problem outside their original field. The classifieds is that infrastructure.

As new fields open continuously, each one posted openly for anyone qualified to apply against, changing careers entirely, more than once, becomes as normal on Plurimo as holding one job for an entire working life has been everywhere else. A person is not locked into the discipline they originally trained in; they are free to move toward whatever emergent problem their particular pattern-recognition actually fits, as often as new problems continue to emerge.

This rests on a specific claim about what humans are for, relative to AI, and the document states it directly rather than leaving it implicit. AI systems, including the domain-specific engines this platform trains, are fundamentally high-powered calculators: excellent at solving known quantities — well-defined problems with a specifiable objective — and poorly suited to a different task entirely: spotting that an unrecognized, not-yet-defined problem exists in the first place, in the gap between two fields nobody has yet connected. That specific capacity — noticing an emergent problem before it has a name, and recognizing which of one's own unrelated skills might apply to it — is a uniquely human, polymathic capacity, not an incidental one. The purpose of Plurimo, at the scale of the whole platform, is to free that capacity at population scale: removing the structural siloing in both education and employment that currently suppresses it in most people, and building the actual infrastructure — the forums, the Knowledge Shepherd role, the classifieds itself — to turn that cross-discipline capacity into real, fundable, paid work.

The Social Forum Scales Up

The modest after-class hangout launched alongside Phase One grows, at this same transition point, into the full generative sandbox described above: still governed by the same two content restrictions and the same effort-based visibility judgment established at launch, now simply open to the wider scope and larger user base the platform has grown into by this phase.

Paid Real-World Engagement: An Earned VIP Tier, Not a Travel Agency

Eligibility for paid real-world conference and play time is tied to demonstrated, already-realized contribution, not to raw time spent on the platform. Two paths qualify a premium user: genuine equity or royalty income already being earned — a master teacher who has trained one of the intelligence engines, or a Knowledge Shepherd already receiving royalties from a field they helped create — or, for premium social-forum users, a sustained record of high-engagement content on the work side of the platform as well as the social side in addition to premium content driving high levels of engagement (past a certain threshold) on the social media side, which substitutes for equity or royalty income as proof of real value being generated for the site.

Eligibility converts to actual conference and play days through a points system, not a flat entitlement: points accrue based on the amount of equity earned (itself a function of training volume completed), the amount of royalty income earned (itself a function of the number of successful fields a Shepherd has helped create), or the rate of engagement a user's own social content is driving. Up to two weeks each of paid work-conference and paid play-event time annually remains the ceiling, usable in smaller increments or as single long blocks.

Accountability and Anti-Gaming Design

Work-conference attendees must upload actual, checkable contributed work to the academic forum; play-conference attendees must upload a play summary to the social forum, with any created game added to a Creative-Commons games forum. Work emerging accidentally during play may become a genuine new academic-forum thread, but only after the event concludes and tagged to its play origin — work is never a permitted, deliberate goal during play itself, which is what prevents either side's engagement from being gamed as the other.

Cheating costs a full year's accumulated days, rebuilt from scratch at a pace that can take years given how much sustained contribution the top tier requires; a second offense results in a permanent ban.

This eligibility structure is a deliberate correction, worth stating directly: tying access to already-realized equity, royalties, or demonstrated engagement value — rather than to raw platform usage — makes the pool of eligible users scale with value the platform has already captured, not with time spent alone. This is what keeps the program financially sound without requiring an external sponsorship model to fund it, and it specifically prevents the failure mode of highly active, low-income users maximizing engagement metrics without having contributed anything the platform has already monetized. The population that qualifies will still be substantial — which is the goal, since it still builds a genuine global Plurimo community that extends beyond the platform itself — but meaningfully smaller and self-funding in a way raw engagement-based eligibility could never be.

One open question remains, named directly rather than smoothed over: the concrete verification mechanism behind "checkable" work and play uploads, which needs to be at least as well-specified as the penalty it supports.

Parallel and Ongoing Growth: Domain Engines and Training Tracks

These build continuously from Phase One onward, in parallel with the phase-transition timeline above, rather than waiting for it.

The Writing Intelligence Engine

Emerging directly from the Plurimo platform, the Writing Intelligence Engine is a domain-aware writing analysis and improvement tool — and eventually a fully trained AI writing expert — built on the richest, most ethically sourced writing pedagogy dataset ever assembled. Existing AI writing tools treat writing as monolithic; they do not understand that the rules of good writing in a legal brief are genuinely different from literary fiction, from scientific papers, or from journalism. Correctness is context-dependent, and no existing tool understands context at the level of discourse community and genre.

The AI comes preloaded with thousands of possible reasons for any given kind of edit, proposing a tag automatically. Right-click always approves the current tag in one motion. Left-click always opens the cell for writing — whether because the AI's proposed tag is wrong and needs

replacing, or because no tag exists yet and one needs to be created from scratch, cursor placed immediately before the change.

A human moderation layer governs what enters the training set: an approved-but-incorrect tag is removed from training; a teacher-written incorrect correction is removed; a correct tag in non-standard language is normalized before entering the database; a correction that fixes a genuine AI error, in non-standard language, is normalized and retained as real AI-error training data, with the paper's own tag updated to match.

This concentrates teacher effort exactly where it has the highest training value, while leaving teachers with something valuable independent of the AI's own training: a persistent, tagged record of why each correction was made, usable in mathematics and history as readily as in an English paper, since the mechanism captures the reasoning behind a change, not the content of the change itself.

Master teachers who train this engine are compensated with pay and equity, not pay alone. Equity vests only upon completion of the full training period and is forfeited, though earned pay is not, for a sustained pattern of bad tagging or bad approvals.

Participation in data collection is entirely opt-in, with both teacher and student required to consent independently, available on the free tier at no charge, fully informed, symmetrical, and non-coercive.

The Expansion Roadmap: A Family of Domain-Specific Intelligence Engines

The Writing Intelligence Engine is the proof of concept, not the ceiling. The same architecture extends to every subject on the platform. The next planned release is a STEM Intelligence Engine, trained the same way on real expert correction data across mathematics and the sciences. Later releases are planned as compiled, broader engines organized by discipline family as sufficient master-teacher participation accumulates in each area — a Humanities Intelligence Engine, an Arts Intelligence Engine, and others as they become viable, alongside continued subject-specific engines where warranted.

Software releases under this roadmap are ongoing and cumulative, not a single fixed milestone: each new cohort of onboarded master teachers in a given subject area moves that subject's own engine closer to release. Once sufficiently trained, each engine becomes a purchasable product — licensable to individual users, school systems, universities, and corporate training programs.

The Site Algorithm Training Track

Free-tier users — students and teachers who could not otherwise afford premium access — may train a distinct AI governing the site's own matching algorithm, particularly for students who have completed the CMTR pedagogical evaluation above. This labor is compensated at the same modest rate paid to Writing Intelligence Engine trainers.

This arrangement is reciprocal and closed-loop by design: the specific communities who contribute the training labor are the communities the resulting software is released back to, once development is complete — not a general promise of eventual free access for everyone.

Free premium-tier access specifically follows genuine, active, interactive labor, not compensation in general. Teachers who train either the Writing Intelligence Engine or the site-matching algorithm receive free premium access, in addition to their pay and equity, because both require real, ongoing engagement with the training task itself. Students who train the CMTR matching engine also receive free premium access, for the same reason — that training genuinely is interactive work. Students whose essays contribute to the Writing Intelligence Engine's training data are different in kind, not merely in scale: their own role is passive consent to having already-completed homework used as training material during a teacher's editing process, not active, ongoing labor. They are compensated for that consent, but do not additionally receive free premium access on top of it, since no interactive work occurred that would justify the same benefit teachers and CMTR-training students actually did the work to earn.

The CMTR Matching Engine

A separate matching engine, trained specifically on the founder's own Cognitive Trait Manifold Ratio data, optimizing teacher-preset, student-preset, and teacher-student matching directly. Unlike the site-algorithm track above, the trained CMTR software itself is intended for free public release to the communities that train it once complete, framed here as part of Plurimo's own stated philanthropic aims, with participants explicitly consenting to that eventual free release as a condition of training it.

Funding Strategy

AI-assisted development tools significantly reduce engineering costs, enabling a functional Phase One MVP with a lean team.

Target Funding Sources

- Grants (non-dilutive): NSF SBIR, Department of Education, Gates Foundation, Chan Zuckerberg Initiative, Simons Foundation, UNESCO and international development bodies.
- Equity investment: edtech-focused VCs (Reach Capital, Learn Capital, Owl Ventures), impact investors, angel investors, Y Combinator.
- Crowdfunding as a simultaneous capital, validation, and community-building mechanism.

Founding Philosophy: Democratized Free Thought

Plurimo is built on a single philosophical commitment: free thought, meaning genuine intellectual plurality, rigorously enforced.

And free thought does sometimes need to be enforced. Left to their own devices, even well-meaning teachers naturally present what they already believe as simply the way things are. The Plurimo classification and flag system is the structural protection that makes genuine intellectual freedom real, just as free speech requires structural legal protection to survive in practice.

This is not ideological neutrality, which is merely a claim anyone can make. It is architectural enforcement of epistemic standards, verifiable in a way no other platform can match. The result is a platform that cannot be captured by any ideological direction, because it does not favor conclusions; it favors rigor.

The Long-Term Vision

Plurimo begins as a tutoring platform. It becomes an independent knowledge economy. It develops a compensation structure that rewards expertise and learning in equal measure. It trains AI on the richest, most ethically sourced pedagogical dataset ever assembled. And it ends as infrastructure for human learning itself, scalable from a single homeschooled child to a national school system.

Every person who teaches on Plurimo makes the platform better for every student. Every student who learns on Plurimo contributes to an AI that will teach the next student better. Every researcher who finds a peer through the forums expands the network of independent intellectual life. Every expert whose knowledge is embedded in a preset trains the next generation of teachers who will never meet them.

This is the independent knowledge economy. This is Plurimo.