

Actor Tracker — Seed Note

Status: queued, not started. Pick up after a few more general (effect-layer) trackers are built.

The Core Distinction This Project Is Built On

Three layers, not two:

1. Mechanism — the actual actions of an identifiable actor (a specific administrator diverting funds, a specific office deciding which grants to cancel and when). This layer is never observed by any tracker built so far.
2. Effect, functioning as general cause — the aggregate, measurable consequence those actions leave behind (a resource trajectory that stays flat while adaptive capacity erodes; a debt-spike hitting several institutions simultaneously; a resource inflow smaller than expected by a specific, unexplained dollar amount). This is what every tracker built so far actually measures, and it is the real, correct causal level for predicting a flip — the flip responds to this layer, not to the underlying actions directly.
3. Flip — the classified outcome itself.

Existing terminology in the framework (extraction flux, push-vectors, the propagation machinery) already lives at layer 2, the effect layer. Naming the new project "actor tracker" rather than "vector tracker" was a deliberate choice to avoid blurring that distinction — "vector" is already doing real work one layer up from where this project lives.

What This Project Actually Is

Not a data-science extension of the existing trackers. A genuinely different methodology: forensic data science — closer in practice to forensic accounting or investigative journalism than to the IPEDS-style structured-data work done so far. The existing trackers work because their data (IPEDS, Form 990s, bond filings) is effect-layer by construction: institutions are required to report outcomes on a regular schedule, regardless of whether anything went wrong.

An actor tracker needs mechanism-layer evidence instead: board minutes, deposition transcripts, whistleblower testimony, subpoenaed correspondence, OIG and state AG investigative findings. This kind of evidence is:

- Adversarial to obtain — no actor voluntarily files a report describing their own predatory action the way an institution files an annual finance survey.
- Fragmentary by nature — it typically exists at all only because something already went wrong and someone investigated after the fact, not on a predictable reporting cycle.
- Scale-appropriate to a research program, not a solo project — Carmen's own estimate, stated directly: this would require a whole graduate program's worth of researchers and investigators to assemble properly, not an individual effort.

What the Actor Tracker Would Actually Be Looking For

Not generic evidence of "a decision was made somewhere." Specifically: evidence of the named mechanisms RICD already formalizes mathematically but that the effect-layer trackers built so far can only ever infer indirectly, through their downstream consequences. Confirmed examples from the manuscript itself:

- Frantic Striking Phase (§8.8b) — the search-and-triage cost pattern of repeated, failing survival campaigns.
- Lateral knocks / vented heat (§8.6a, $H_{i,j}$) — an unmetabolized failure discharged onto a neighboring container.
- Voluntary Mutual Sacrifice (§10.5b.11b.iii, "sacrificial lifeboat" in working shorthand) — two comparably-struggling containers where the discharge asymmetry resolves in favor of one's survival through the other's own trajectory redirection, distinct from predatory targeting (§10.5b.11b.i) or accidental capture via declining connectivity (§10.5b.11b.ii).

The higher-ed tracker and its kind — call them collapse-causal layer trackers, or equivalently downstream consequence trackers, per Carmen's naming — measure the aggregate effect these mechanisms leave behind (a resource trajectory, a debt-spike, an extraction's dollar figure) without ever needing to identify which named mechanism, specifically, produced it. An actor tracker's whole purpose is the piece that leaves out: given a real, observed collapse, was it actually Frantic Striking, a lateral knock, a sacrificial-lifeboat trade, predatory targeting, or accidental capture — a real, falsifiable attribution question the effect-layer data was never built to answer, requiring exactly the mechanism-layer evidence described above.

What "Done" Would Look Like, Eventually

Two deliverables, matching the pattern already used for the higher-ed tracker:

1. A prospective adapter — real fields, real sourcing requirements, scoped to what forensic-layer evidence can actually supply (case files, not CSVs).
2. A how-to manual — but a genuinely different kind of manual than the existing one. Not "how to extract field F2A18 from a CSV." Closer to: how to build a case file: what counts as evidence of a specific decision versus merely being suggestive of one, how to establish provenance and reliability for fragmentary, adversarially-obtained sources, and how to avoid quietly sliding a layer-2 (effect) finding into a layer-1 (mechanism) claim it doesn't actually support.

A Real, Concrete Case Worth Using as a Working Example

Abundant Blessings, a Hyde Park charity: its executive director was charged in a real, documented \$23 million federal fraud case (HUD-OIG and FBI investigation), built on fabricated vendor invoices using stolen real company names and logos, and self-dealing rent payments disguised as market-rate third-party leases. The detail that matters most for this project specifically: when a LAHSA investigator asked whether the charity's board was aware of how money was being spent, it turned out the board itself was fabricated — some listed members did not exist, others had never heard of the organization or its director.

This surfaces a real complication the higher-ed tracker never had to face. External Governance Attestation (RICD Definition 6.1.1) worked in that project because the attesting bodies — accreditors, AAUP — are genuinely independent of the institution being assessed, and their findings are a matter of public record regardless of what the institution itself says. An actor tracker cannot assume this. In a case like Abundant Blessings, the nominal oversight structure is itself part of the fabrication, not a real, independent check on it. Any actor-tracker adapter would need its own way of verifying that a claimed governance or oversight structure actually exists and is actually independent, before treating its attestations as trustworthy — a genuinely new kind of validation problem, not a variant of anything the effect-layer trackers needed.

A sharper way to state this, added after later conversation (RICD 10.2, flagged there as conceptual, not yet tested against data): a container's real boundary is constituted by wherever actor energy is actually interfacing with a recognizable pattern — its energetic signature — not by any official claim about that boundary. Ordinary institutional decline is the signature drifting from real activity slowly enough that the gap eventually leaks into a container's own downstream numbers, which is what let this project's effect-layer tracker catch real decline years early without ever observing an actor directly. Fraud is different in kind, not degree: it is the manufacture of a signature specifically engineered to have no real container behind it, built so an outside observer mistakes the claim for the fact — Abundant Blessings' fabricated board is a clean instance. The working hypothesis for why an actor tracker specifically needs mechanism-layer evidence, stated here as something to test rather than something already confirmed: ordinary drift leaks into effect-layer data on its own, but a signature engineered to fake its own container is built to suppress exactly that leak, which is plausibly why effect-layer data alone will not be enough once fraud, rather than ordinary decline, is the actual phenomenon.

Open Question for Later, Not Now

Whether an actor tracker's classifier target is even the same kind of thing as a collapse/flip classification, or whether it's structurally a different task entirely (e.g., attributing a known effect to a specific actor, rather than predicting an

outcome from features). Worth thinking through once real candidate case material is in hand, not from first principles now.